

The background of the slide is a light gray gradient. It is decorated with numerous realistic water droplets of various sizes. Some droplets are large and prominent, while others are small and subtle. They are scattered across the slide, with a higher concentration in the top-left and bottom-right corners. The droplets have highlights and shadows, giving them a three-dimensional appearance.

# TRANSACTION MANAGEMENT

RELATIONAL DATABASES

# TRANSACTIONS

- TRANSACTIONS ARE A FUNDAMENTAL CONCEPT OF ALL DATABASE SYSTEMS
- TRANSACTIONS ALLOW USERS TO MAKE CHANGES TO DATA AND THEN DECIDE WHETHER TO SAVE OR DISCARD THE WORK
- DATABASE TRANSACTIONS BUNDLE MULTIPLE STEPS INTO ONE LOGICAL UNIT OF WORK

# TRANSACTIONS

- A TRANSACTION CONSISTS OF ONE OF THE FOLLOWING:
  - DML STATEMENT WHICH CONSTITUTE ONE CONSISTENT CHANGE TO THE DATA
  - INSERT, UPDATE, DELETE AND MERGE
  - ONE DDL STATEMENT SUCH AS CREATE, ALTER, DROP, RENAME, OR TRUNCATE
  - ONE DCL STATEMENT SUCH AS GRANT OR REVOKE

# TRANSACTIONS

- A BANK DATABASE CONTAINS BALANCES FOR VARIOUS CUSTOMER ACCOUNTS, AS WELL AS DEPOSIT BALANCES FOR OTHER BRANCHES
- A CUSTOMER WANTS TO WITHDRAW AND TRANSFER MONEY FROM HIS ACCOUNT AND DEPOSIT IT INTO ANOTHER CUSTOMER'S ACCOUNT AT A DIFFERENT BRANCH
- SEVERAL STEPS INVOLVED, BOTH BRANCHES WANT TO ENSURE ALL STEPS HAPPEN, OR NONE OF THEM HAPPEN AND IF THE SYSTEM CRASHES THE TRANSACTION IS NOT PARTIALLY COMPLETE.
- GROUPING THE WITHDRAWAL AND DEPOSIT STEPS INTO ONE TRANSACTION PROVIDES THIS GUARANTEE

# TRANSACTIONS

- TRANSACTIONS ARE CONTROLLED USING THE FOLLOWING STATEMENTS:
  - COMMIT: REPRESENTS THE POINT IN TIME WHERE THE USER HAS MADE ALL THE CHANGES HE/SHE WANTS TO HAVE LOGICALLY GROUPED TOGETHER, AND BECAUSE NO MISTAKES HAVE BEEN MADE, THE USER IS READY TO SAVE THE WORK. WHEN A COMMIT STATEMENT IS ISSUED THE CURRENT TRANSACTION ENDS, MAKING ALL PENDING CHANGES PERMANENT.
  - ROLLBACK: ENABLES THE USER TO DISCARD ALL PENDING CHANGES MADE TO THE DATABASE
  - SAVEPOINT: CREATES A MARKER IN A TRANSACTION, WHICH DIVIDES IT INTO SMALLER PIECES.
  - ROLLBACK TO SAVEPOINT: ALLOWS THE USER TO ROLLBACK TO A SPECIFIED SAVEPOINT, THUS DISCARDING ONLY THOSE CHANGES MADE AFTER THE SAVEPOINT

# TRANSACTIONS

- A TRANSACTION BEGINS WITH THE FIRST DML STATEMENT
  - IT ENDS WHEN ONE OF THE FOLLOWING OCCURS:
    - A COMMIT OR ROLLBACK STATEMENT IS ISSUED
    - A DDL STATEMENT IS ISSUED
    - A DCL STATEMENT IS ISSUED
    - A USER EXITS NORMALLY FROM THE DATABASE UTILITY, CAUSING THE CURRENT TRANSACTION TO BE IMPLICITLY COMMITTED.
- A DDL STATEMENT OR DCL STATEMENT IS AUTOMATICALLY COMMITTED THEREFORE IMPLICITLY ENDS A TRANSACTION
- EVERY DATA CHANGE MADE DURING A TRANSACTION IS TEMPORARY UNTIL THE TRANSACTION IS COMMITTED.

# READ CONSISTENCY

- READ CONSISTENCY GUARANTEES A CONSISTENT VIEW OF THE DATA BY ALL USERS AT ALL TIMES
- READERS DO NOT VIEW DATA THAT IS IN THE PROCESS OF BEING CHANGED
- WRITERS ARE ENSURED THAT THE CHANGES TO THE DATABASE ARE DONE IN A CONSISTENT WAY
- CHANGES MADE BY ONE WRITER DO NOT DESTROY OR CONFLICT WITH CHANGES ANOTHER WRITER IS MAKING

# CONCURRENCY

- READ CONSISTENCY IS AN AUTOMATIC IMPLEMENTATION
- THIS IS TO ENSURE THAT IN A MULTI USER ENVIRONMENT SUCH AS RELATIONAL DATABASE THAT TRANSACTIONS THAT ARE RUNNING AT THE SAME TIME WORKING ON THE SAME PIECES OF DATA DO NOT INTERFERE WITH EACHOTHER.



# CONCURRENCY

- A PARTIAL COPY OF THE DATABASE IS KEPT IN UNDO SEGMENTS.
- WHEN USER A ISSUES AN INSERT, UPDATE OR DELETE OPERATION TO THE DATABASE, THE ORACLE SERVER TAKES A SNAPSHOT (COPY) OF THE DATA BEFORE IT IS CHANGED AND WRITES IT TO THE UNDO (ROLLBACK) SEGMENT.
- USER B STILL SEES THE DATABASE AS IT EXISTED BEFORE THE CHANGES STARTED; HE VIEWS THE UNDO SEGMENT'S SNAPSHOT OF THE DATA.
- BEFORE CHANGES ARE COMMITTED TO THE DATABASE, ONLY THE USER MAKING THE CHANGE SEES THEM, EVERYONE ELSE SEES THE SNAPSHOT.
- THIS GUARANTEES THAT READERS OF THE DATA SEE CONSISTENT DATA THAT IS NOT CURRENTLY UNDERGOING CHANGE.

# CONCURRENCY

- WHEN A DML IS COMMITTED, THE CHANGE MADE TO THE DATABASE BECOMES VISIBLE TO ANYONE EXECUTING A SELECT STATEMENT.
- IF THE TRANSACTION IS ROLLED BACK, THE CHANGES ARE UNDONE:
  - THE ORIGINAL, OLDER VERSION OF DATA IN THE UNDO SEGMENT IS WRITTEN BACK TO THE TABLE
  - ALL USERS SEE THE DATABASE AS IT EXISTED BEFORE THE TRANSACTION BEGAN.

# CONCURRENCY PROBLEMS

- **LOST UPDATE PROBLEM:**
  - WHEN TWO TRANSACTIONS ARE UPDATING THE SAME DATA ELEMENT AND ONE OF THE UPDATES IS LOST (OVERWRITTEN BY THE OTHER TRANSACTION)
  - IF TRANSACTION A READS THE DATA VALUE BEFORE TRANSACTION B COMMITS THEN THE CHANGE TRANSACTION B MAKES IS LOST

| Time | Transaction | Step            | Stored value     |
|------|-------------|-----------------|------------------|
| 1    | T1          | Read prod_qty   | 35               |
| 2    | T2          | Read prod_qty   | 35               |
| 3    | T1          | Prod_qty=35+100 | 135              |
| 4    | T2          | Prod_qty=35-30  | 5                |
| 5    | T1          | Write(prod_qty) | 135 (lost update |
| 6    | T2          | Write(prod_qty) | 5                |

# CONCURRENCY PROBLEMS

- UNCOMMITTED DATA:
  - THIS OCCURS WHEN TWO TRANSACTIONS ARE EXECUTED AT THE SAME TIME AND ONE ROLLS BACK AFTER THE OTHER TRANSACTION ALREADY ACCESSED THE UNCOMMITTED DATA.

| Time | Transaction | Step            | Stored value                           |
|------|-------------|-----------------|--|
| 1    | T1          | Read prod_qty   | 35                                     |
| 2    | T1          | Prod_qty=35+100 |  |
| 3    | T1          | Write(prod_qty) | 135                                    |
| 4    | T2          | Read prod_qty   | 135 <sub>(read uncommitted data)</sub> |
| 5    | T2          | Prod_qty=135-30 | 105                                    |
| 6    | T1          | Rollback        | 35                                     |
| 7    | T2          | Write(prod_qty) | 105                                    |

# CONCURRENCY PROBLEMS

- INCONSISTENT RETRIEVALS:
  - THIS OCCURS WHEN A TRANSACTION ACCESSES DATA BEFORE AND AFTER ANOTHER TRANSACTION(S) FINISHES WORKING WITH DATA.
  - T1 IS KEEPING A RUNNING TOTAL OF PRODUCT QUANTITIES
  - T2 IS ADDING TO THE QUANTITIES OF PRODUCTS

| Time | Transaction | Step                                 | Stored value | Sum              |
|------|-------------|--------------------------------------|--------------|------------------|
| 1    | T1          | Read prod_qty for id='73FG'          | 35           | 35               |
| 2    | T1          | Read prod_qty for id='23TY'          | 55           | 90               |
| 3    | T2          | Read prod_qty for id='11BA'          | 10           |                  |
| 4    | T2          | Prod_qty('11BA')=10+15               | 25           |                  |
| 5    | T2          | Write(prod_qty for id='11BA')        | 25           |                  |
| 6    | T1          | Read prod_qty for id='11BA'          | 25           | 115(read after)  |
| 7    | T1          | Read prod_qty for id='55KK'          | 50           | 165(read before) |
| 8    | T2          | Read prod_qty for id='55KK'          | 50           |                  |
| 9    | T2          | Prod_qty('55KK')=50+20               | 70           |                  |
| 10   | T2          | Write(prod_qty for id='55KK') Commit | 70           |                  |
| 11   | T1          | Read prod_qty for id='99UH'          | 5            | 170              |

# CONCURRENCY CONTROL TECHNIQUES

- ▶ TWO BASIC CONCURRENCY CONTROL TECHNIQUES ARE:
  - ▶ LOCKING
  - ▶ TIMESTAMPING
  
- ▶ BOTH ARE CONSERVATIVE (PESSIMISTIC) APPROACHES: DELAY TRANSACTIONS IN CASE THEY CONFLICT WITH OTHER TRANSACTIONS.
  
- ▶ OPTIMISTIC METHODS ASSUME CONFLICT IS RARE AND ONLY CHECK FOR CONFLICTS AT COMMIT.
  - ▶ VERSIONING

# LOCKING

- TRANSACTION USES LOCKS TO DENY ACCESS TO OTHER TRANSACTIONS AND SO PREVENT INCORRECT UPDATES.
- MOST WIDELY USED APPROACH TO ENSURE READ CONSISTENCY.
- GENERALLY, A TRANSACTION MUST CLAIM A SHARED (READ) OR EXCLUSIVE (WRITE) LOCK ON A DATA ITEM BEFORE READ OR WRITE.
- LOCK PREVENTS ANOTHER TRANSACTION FROM MODIFYING ITEM OR EVEN READING IT, IN THE CASE OF A WRITE LOCK.

# LOCKING - BASIC RULES

- IF TRANSACTION HAS **SHARED** LOCK ON ITEM, IT CAN READ BUT NOT UPDATE ITEM.
- IF TRANSACTION HAS **EXCLUSIVE** LOCK ON ITEM, IT CAN BOTH READ AND UPDATE ITEM.
- READS CANNOT CONFLICT, SO MORE THAN ONE TRANSACTION CAN HOLD SHARED LOCKS SIMULTANEOUSLY ON SAME ITEM.
- EXCLUSIVE LOCK GIVES TRANSACTION EXCLUSIVE ACCESS TO THAT ITEM (I.E. THE ITEM CANNOT EVEN BE READ BY OTHER TRANSACTIONS).



# LOCKING - BASIC RULES

- SOME SYSTEMS ALLOW TRANSACTION TO UPGRADE READ LOCK TO AN EXCLUSIVE LOCK, OR DOWNGRADE EXCLUSIVE LOCK TO A SHARED LOCK.

# LOCKING

- LOCKING DOES NOT GUARANTEE READ CONSISTENCY
- IF A TRANSACTION RELEASES THE LOCK PRIOR TO A COMMIT THEN ANOTHER TRANSACTION COULD READ THE DATA AND THE ORIGINAL TRANSACTION END UP ROLLING BACK RATHER THAN COMMITTING.
- UNCOMMITTED DATA PROBLEM

# TWO-PHASE LOCKING (2PL)

- TRANSACTION FOLLOWS 2PL PROTOCOL IF ALL LOCKS ARE HELD UNTIL THE TRANSACTION ENDS (EITHER COMMITTED OR ROLLED BACK)
- TWO PHASES FOR TRANSACTION:
  - GROWING PHASE - ACQUIRES ALL LOCKS BUT CANNOT RELEASE ANY LOCKS.
  - SHRINKING PHASE - RELEASES LOCKS BUT CANNOT ACQUIRE ANY NEW LOCKS.

# PREVENTING LOST UPDATE PROBLEM USING 2PL

| Time            | T <sub>1</sub>                                       | T <sub>2</sub>  | bal <sub>x</sub> |
|-----------------|--|---|------------------|
| t <sub>1</sub>  |  | begin_transaction                                     | 100              |
| t <sub>2</sub>  | begin_transaction                                    | write_lock( <b>bal<sub>x</sub></b> )                  | 100              |
| t <sub>3</sub>  | write_lock( <b>bal<sub>x</sub></b> )                 | read( <b>bal<sub>x</sub></b> )                        | 100              |
| t <sub>4</sub>  | WAIT   | <b>bal<sub>x</sub></b> = <b>bal<sub>x</sub></b> + 100 | 100              |
| t <sub>5</sub>  | WAIT   | write( <b>bal<sub>x</sub></b> )                       | 200              |
| t <sub>6</sub>  | WAIT   | commit/unlock( <b>bal<sub>x</sub></b> )               | 200              |
| t <sub>7</sub>  | read( <b>bal<sub>x</sub></b> )                       |   | 200              |
| t <sub>8</sub>  | <b>bal<sub>x</sub></b> = <b>bal<sub>x</sub></b> - 10 |   | 200              |
| t <sub>9</sub>  | write( <b>bal<sub>x</sub></b> )                      |   | 190              |
| t <sub>10</sub> | commit/unlock( <b>bal<sub>x</sub></b> )              |   | 190              |

# PREVENTING UNCOMMITTED DATA PROBLEM USING 2PL

| Time            | T <sub>3</sub>                                       | T <sub>4</sub>  | bal <sub>x</sub> |
|-----------------|--|---|------------------|
| t <sub>1</sub>  |  | begin_transaction                                     | 100              |
| t <sub>2</sub>  |  | write_lock( <b>bal<sub>x</sub></b> )                  | 100              |
| t <sub>3</sub>  |  | read( <b>bal<sub>x</sub></b> )                        | 100              |
| t <sub>4</sub>  | begin_transaction                                    | <b>bal<sub>x</sub></b> = <b>bal<sub>x</sub></b> + 100 | 100              |
| t <sub>5</sub>  | write_lock( <b>bal<sub>x</sub></b> )                 | write( <b>bal<sub>x</sub></b> )                       | 200              |
| t <sub>6</sub>  | WAIT   | rollback/unlock( <b>bal<sub>x</sub></b> )             | 100              |
| t <sub>7</sub>  | read( <b>bal<sub>x</sub></b> )                       |   | 100              |
| t <sub>8</sub>  | <b>bal<sub>x</sub></b> = <b>bal<sub>x</sub></b> - 10 |   | 100              |
| t <sub>9</sub>  | write( <b>bal<sub>x</sub></b> )                      |   | 90               |
| t <sub>10</sub> | commit/unlock( <b>bal<sub>x</sub></b> )              |   | 90               |

# PREVENTING INCONSISTENT ANALYSIS

| Time            | T <sub>5</sub>   | T <sub>6</sub>  | bal <sub>x</sub> | bal <sub>y</sub> | bal <sub>z</sub> | sum |
|-----------------|--|---|------------------|------------------|------------------|-----|
| t <sub>1</sub>  |  | begin_transaction   | 100              | 50               | 25               |     |
| t <sub>2</sub>  | begin_transaction  | sum = 0   | 100              | 50               | 25               | 0   |
| t <sub>3</sub>  | write_lock( <b>bal<sub>x</sub></b> )                     |   | 100              | 50               | 25               | 0   |
| t <sub>4</sub>  | read( <b>bal<sub>x</sub></b> )                           | read_lock( <b>bal<sub>x</sub></b> )                                       | 100              | 50               | 25               | 0   |
| t <sub>5</sub>  | <b>bal<sub>x</sub> = bal<sub>x</sub> - 10</b>            | WAIT  | 100              | 50               | 25               | 0   |
| t <sub>6</sub>  | write( <b>bal<sub>x</sub></b> )                          | WAIT  | 90               | 50               | 25               | 0   |
| t <sub>7</sub>  | write_lock( <b>bal<sub>z</sub></b> )                     | WAIT  | 90               | 50               | 25               | 0   |
| t <sub>8</sub>  | read( <b>bal<sub>z</sub></b> )                           | WAIT  | 90               | 50               | 25               | 0   |
| t <sub>9</sub>  | <b>bal<sub>z</sub> = bal<sub>z</sub> + 10</b>            | WAIT  | 90               | 50               | 25               | 0   |
| t <sub>10</sub> | write( <b>bal<sub>z</sub></b> )                          | WAIT  | 90               | 50               | 35               | 0   |
| t <sub>11</sub> | commit/unlock( <b>bal<sub>x</sub>, bal<sub>z</sub></b> ) | WAIT  | 90               | 50               | 35               | 0   |
| t <sub>12</sub> |  | read( <b>bal<sub>x</sub></b> )  | 90               | 50               | 35               | 0   |
| t <sub>13</sub> |  | sum = sum + <b>bal<sub>x</sub></b>  | 90               | 50               | 35               | 90  |
| t <sub>14</sub> |  | read_lock( <b>bal<sub>y</sub></b> )                                       | 90               | 50               | 35               | 90  |
| t <sub>15</sub> |  | read( <b>bal<sub>y</sub></b> )  | 90               | 50               | 35               | 90  |
| t <sub>16</sub> |  | sum = sum + <b>bal<sub>y</sub></b>  | 90               | 50               | 35               | 140 |
| t <sub>17</sub> |  | read_lock( <b>bal<sub>z</sub></b> )                                       | 90               | 50               | 35               | 140 |
| t <sub>18</sub> |  | read( <b>bal<sub>z</sub></b> )  | 90               | 50               | 35               | 140 |
| t <sub>19</sub> |  | sum = sum + <b>bal<sub>z</sub></b>  | 90               | 50               | 35               | 175 |
| t <sub>20</sub> |  | commit/unlock( <b>bal<sub>x</sub>, bal<sub>y</sub>, bal<sub>z</sub></b> ) | 90               | 50               | 35               | 175 |

# DEADLOCK

- AN IMPASSE THAT MAY RESULT WHEN TWO (OR MORE) TRANSACTIONS ARE EACH WAITING FOR LOCKS HELD BY THE OTHER TO BE RELEASED.

| Time            | T <sub>17</sub>                                      | T <sub>18</sub>                                       |
|-----------------|--|---|
| t <sub>1</sub>  | begin_transaction                                    |   |
| t <sub>2</sub>  | write_lock( <b>bal<sub>x</sub></b> )                 | begin_transaction                                     |
| t <sub>3</sub>  | read( <b>bal<sub>x</sub></b> )                       | write_lock( <b>bal<sub>y</sub></b> )                  |
| t <sub>4</sub>  | <b>bal<sub>x</sub></b> = <b>bal<sub>x</sub></b> - 10 | read( <b>bal<sub>y</sub></b> )                        |
| t <sub>5</sub>  | write( <b>bal<sub>x</sub></b> )                      | <b>bal<sub>y</sub></b> = <b>bal<sub>y</sub></b> + 100 |
| t <sub>6</sub>  | write_lock( <b>bal<sub>y</sub></b> )                 | write( <b>bal<sub>y</sub></b> )                       |
| t <sub>7</sub>  | WAIT   | write_lock( <b>bal<sub>x</sub></b> )                  |
| t <sub>8</sub>  | WAIT   | WAIT  |
| t <sub>9</sub>  | WAIT   | WAIT  |
| t <sub>10</sub> | ⋮  | WAIT  |
| t <sub>11</sub> | ⋮  | ⋮   |

# DEADLOCK

- ONLY ONE WAY TO BREAK DEADLOCK: ABORT ONE OR MORE OF THE TRANSACTIONS.
- DEADLOCK SHOULD BE TRANSPARENT TO USER, SO DBMS SHOULD RESTART TRANSACTION(S).
- THREE GENERAL TECHNIQUES FOR HANDLING DEADLOCK:
  - TIMEOUTS.
  - DEADLOCK PREVENTION.
  - DEADLOCK DETECTION AND RECOVERY.



# TIMEOUTS

- TRANSACTION THAT REQUESTS LOCK WILL ONLY WAIT FOR A SYSTEM-DEFINED PERIOD OF TIME.
- IF LOCK HAS NOT BEEN GRANTED WITHIN THIS PERIOD, LOCK REQUEST TIMES OUT.
- IN THIS CASE, DBMS ASSUMES TRANSACTION MAY BE DEADLOCKED, EVEN THOUGH IT MAY NOT BE, AND IT ABORTS AND AUTOMATICALLY RESTARTS THE TRANSACTION.


# DEADLOCK PREVENTION

- DBMS LOOKS AHEAD TO SEE IF TRANSACTION WOULD CAUSE DEADLOCK AND NEVER ALLOWS DEADLOCK TO OCCUR.
- COULD ORDER TRANSACTIONS USING TRANSACTION TIMESTAMPS:
  - WAIT-DIE: ONLY AN OLDER TRANSACTION CAN WAIT FOR YOUNGER ONE, OTHERWISE TRANSACTION IS ABORTED (DIES) AND RESTARTED WITH SAME TIMESTAMP.
  - WOUND-WAIT: ONLY A YOUNGER TRANSACTION CAN WAIT FOR AN OLDER ONE. IF OLDER TRANSACTION REQUESTS LOCK HELD BY YOUNGER ONE, YOUNGER ONE IS ABORTED (WOUNDED).

# DEADLOCK DETECTION AND RECOVERY

- DBMS ALLOWS DEADLOCK TO OCCUR BUT RECOGNIZES IT AND BREAKS IT.
- USUALLY HANDLED BY CONSTRUCTION OF WAIT-FOR GRAPH (WFG) SHOWING TRANSACTION DEPENDENCIES:
  - CREATE A NODE FOR EACH TRANSACTION.
  - CREATE EDGE  $T_i \rightarrow T_j$ , IF  $T_i$  WAITING TO LOCK ITEM LOCKED BY  $T_j$ .
- DEADLOCK EXISTS IF AND ONLY IF WFG CONTAINS CYCLE.
- WFG IS CREATED AT REGULAR INTERVALS.

## EXAMPLE - WAIT-FOR-GRAPH (WFG)



# RECOVERY FROM DEADLOCK DETECTION

- SEVERAL ISSUES:
  - CHOICE OF DEADLOCK VICTIM;
  - HOW FAR TO ROLL A TRANSACTION BACK;
  - AVOIDING STARVATION.

# TIMESTAMPING

- TRANSACTIONS ORDERED GLOBALLY SO THAT OLDER TRANSACTIONS, TRANSACTIONS WITH SMALLER TIMESTAMPS, GET PRIORITY IN THE EVENT OF CONFLICT.
- CONFLICT IS RESOLVED BY ROLLING BACK AND RESTARTING TRANSACTION.
- NO LOCKS SO NO DEADLOCK.

# TIMESTAMPING

- **TIMESTAMP**
  - A UNIQUE IDENTIFIER CREATED BY DBMS THAT INDICATES RELATIVE STARTING TIME OF A TRANSACTION.
- CAN BE GENERATED BY USING SYSTEM CLOCK AT TIME TRANSACTION STARTED, OR BY INCREMENTING A LOGICAL COUNTER EVERY TIME A NEW TRANSACTION STARTS.

# TIMESTAMPING

- READ/WRITE PROCEEDS ONLY IF LAST READ/WRITE ON THAT DATA ITEM WAS CARRIED OUT BY AN OLDER TRANSACTION.
- OTHERWISE, TRANSACTION REQUESTING READ/WRITE IS RESTARTED AND GIVEN A NEW TIMESTAMP.
- ALSO TIMESTAMPS FOR DATA ITEMS:
  - READ-TIMESTAMP - TIMESTAMP OF LAST TRANSACTION TO READ ITEM;
  - WRITE-TIMESTAMP - TIMESTAMP OF LAST TRANSACTION TO WRITE ITEM.



# TIMESTAMPING EXAMPLE

| John                                     | Time | Marsha                                   | Bal <sub>x</sub> | Data Item Timestamp          |
|--|------|--|------------------|------------------------------|
| read (bal <sub>x</sub> )                 | t1   |  | 100              | Last read by John at t1      |
|  | t2   | read (bal <sub>x</sub> )                 | 100              | Last read by Marsha at t2    |
| bal <sub>x</sub> = bal <sub>x</sub> - 50 | t3   |  | 50               |                              |
| write (bal <sub>x</sub> ) *              | t4   |  | 50               |                              |
| roll back                                | t5   | bal <sub>x</sub> = bal <sub>x</sub> - 10 | 90               |                              |
|  | t6   | write (bal <sub>x</sub> )                | 90               | Last updated by Marsha at t6 |
| read (bal <sub>x</sub> )                 | t7   |  | 90               | Last read by John at t7      |
|  | t8   |  | 90               |                              |
| bal <sub>x</sub> = bal <sub>x</sub> - 50 | t9   |  | 40               |                              |
| write (bal <sub>x</sub> ) **             | t10  |  | 40               | Last update by John at t10   |
|  | t11  |  |                  |                              |
|  | t12  |  |                  |                              |

# TIMESTAMPING EXAMPLE

\* PROBLEM OCCURS HERE – JOHN HAS TRIED TO UPDATE A DATA ITEM WHICH WAS LAST READ BY ANOTHER TRANSACTION (MARSHA). THEREFORE HIS TRANSACTION MUST BE ROLLED BACK - ABORTED AND RESTARTED, AND GIVEN A NEW TIMESTAMP. **FROM TIME T6 AND ONWARDS MARSHA IS NOW THE OLDER TRANSACTION, WHEREAS JOHN'S TRANSACTION IS THE NEWER OF THE TWO (HAVING BEEN GIVEN A NEW TIMESTAMP).**

\*\* THIS TIME JOHN'S UPDATE TO THE BALANCE IS ALLOWED TO PROCEED, AS THE DATA ITEM HAS NOT BEEN READ/UPDATED BY ANYONE ELSE SINCE HIS TRANSACTION WAS RESTARTED.

# OPTIMISTIC TECHNIQUES

- BASED ON ASSUMPTION THAT CONFLICT IS RARE AND MORE EFFICIENT TO LET TRANSACTIONS PROCEED WITHOUT DELAYS TO ENSURE SERIALIZABILITY.
- AT COMMIT, CHECK IS MADE TO DETERMINE WHETHER CONFLICT HAS OCCURRED.
- IF THERE IS A CONFLICT, TRANSACTION MUST BE ROLLED BACK AND RESTARTED.
- POTENTIALLY ALLOWS GREATER CONCURRENCY THAN TRADITIONAL PROTOCOLS.

# VERSIONING

- VERSIONING OF DATA CAN BE USED TO INCREASE CONCURRENCY.
- BASIC TIMESTAMP ORDERING PROTOCOL ASSUMES ONLY ONE VERSION OF DATA ITEM EXISTS, AND SO ONLY ONE TRANSACTION CAN ACCESS DATA ITEM AT A TIME.
- CAN ALLOW MULTIPLE TRANSACTIONS TO READ AND WRITE DIFFERENT VERSIONS OF SAME DATA ITEM.
- IN MULTIVERSION CONCURRENCY CONTROL, EACH WRITE OPERATION CREATES NEW VERSION OF DATA ITEM WHILE RETAINING OLD VERSION.
- NEW VERSIONS ARE LATER MERGED INTO THE DATABASE; CONFLICTS ARE DEALT WITH IF THEY ARISE

# GRANULARITY OF DATA ITEMS

- SIZE OF DATA ITEMS CHOSEN AS UNIT OF PROTECTION BY CONCURRENCY CONTROL PROTOCOL.
- RANGING FROM COARSE TO FINE:
  - THE ENTIRE DATABASE.
  - A FILE.
  - A PAGE (OR AREA OR DATABASE SPACED).
  - A RECORD.
  - A FIELD VALUE OF A RECORD.

# GRANULARITY OF DATA ITEMS

- TRADEOFF:
  - COARSER, THE LOWER THE DEGREE OF CONCURRENCY;
  - FINER, MORE LOCKING INFORMATION THAT IS NEEDED TO BE STORED.
- BEST ITEM SIZE DEPENDS ON THE TYPES OF TRANSACTIONS.